

Chapter exercise3 from ISLR book

Exercise 3.1

" The null hypotheses to Intercept is that means Except the benefit from the TV, radio and newspaper , other benefit – other budget is larger than 0, which rejects the H_0 which is cost = budget in more than 99%.

The null hypotheses to TV is that the benefit of TV ads possibility not equal than the budget of TV ads. There exists more than 99% possibility that benefit of TV ads is not equal than cost of it.

The null hypotheses to Radio is that the benefit of radio ads possibility not equal than the budget of radio ads. There exists more than 99% possibility that benefit of radio ads is not equal than cost of it.

The null hypotheses to newspaper is that the benefit of newspaper ads possibility not equal than the budget of newspaper ads. There exists more than 99% possibility that benefit of newspaper ads is not equal than cost of it. "

Exercise 3.2

KNN regression averages the closest observations to estimate prediction

KNN classifier assigns classification group based on majority of closest observations.

Exercise 3.5

In [2]:

```
from IPython.display import Latex
Latex(r"""
\begin{eqnarray}
\hat{y}_i = x_i \times \frac{\sum_{i'=1}^n \left( x_{i'} y_{i'} \right)}{\sum_{j=1}^n x_j^2} \quad \hat{y}_i = \sum_{i'=1}^n \frac{\left( x_{i'} y_{i'} \right)}{\sum_{j=1}^n x_j^2} \quad \hat{y}_i = \sum_{i'=1}^n \left( \frac{x_{i'} x_{i'}}{\sum_{j=1}^n x_j^2} \right) \times y_{i'} \quad a_{i'} = \frac{x_{i'} x_{i'}}{\sum_{j=1}^n x_j^2} \quad \end{eqnarray}
""")
```

Out[2]:

```
\begin{eqnarray} \hat{y}_i = x_i \times \frac{\sum_{i'=1}^n \left( x_{i'} y_{i'} \right)}{\sum_{j=1}^n x_j^2} \quad \hat{y}_i = \sum_{i'=1}^n \frac{\left( x_{i'} y_{i'} \right)}{\sum_{j=1}^n x_j^2} \quad \hat{y}_i = \sum_{i'=1}^n \left( \frac{x_{i'} x_{i'}}{\sum_{j=1}^n x_j^2} \right) \times y_{i'} \quad a_{i'} = \frac{x_{i'} x_{i'}}{\sum_{j=1}^n x_j^2} \quad \end{eqnarray}
```

Exercise 3.6

In [3]:

```
from IPython.display import Latex
Latex(r"""
Using \; equation \;(3.4), when \;  $x_i = \bar{x}$ , \; then \;  $\hat{\beta}_1 = 0$  \; and \;  $\hat{\beta}_0 = \bar{y}$  \; and \; the \; equation \; for \;  $\hat{y}_i$  \; evaluates \; to \; equal \;  $\bar{y}$ 
\end{eqnarray}
""")
```

Out[3]:

```
\begin{eqnarray} Using \; equation \;(3.4), when \;  $x_i = \bar{x}$ , \; then \;  $\hat{\beta}_1 = 0$  \; and \;  $\hat{\beta}_0 = \bar{y}$  \; and \; the \; equation \; for \;  $\hat{y}_i$  \; evaluates \; to \; equal \;  $\bar{y}$  \end{eqnarray}
```

Exercise 3.11

In [4]:

```
# part (a)
import numpy as np

np.random.seed(11)

s = np.random.normal(0,1,100)
x = s
d = np.random.normal(0,1,100)

y = 2*x + d

import statsmodels.api as sm

model1 = sm.OLS(y,x).fit()

model1.summary()
```

Out [4]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.790			
Model:	OLS	Adj. R-squared:	0.787			
Method:	Least Squares	F-statistic:	371.4			
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	2.82e-35			
Time:	14:16:24	Log-Likelihood:	-136.94			
No. Observations:	100	AIC:	275.9			
Df Residuals:	99	BIC:	278.5			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t P> t [0.025 0.975]			
x1	1.9754	0.103	19.272	0.000	1.772	2.179
Omnibus:	0.555	Durbin-Watson:	1.775			
Prob(Omnibus):	0.758	Jarque-Bera (JB):	0.687			
Skew:	0.077	Prob(JB):	0.709			
Kurtosis:	2.624	Cond. No.	1.00			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The estimated coefficient 2.04 means with one unit increase in x, y will increase 2.04 units. Standard error rate shows the one standard deviation of beta_x is nearly 0. So the 95% confidence interval of beta_x is [1.872,2.211] P-value close to 0 shows that x is statistically significant.

In [5]:

```
# part (b)

model2 = sm.OLS(x,y).fit()

model2.summary()
```

Out [5]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.790
Model:	OLS	Adj. R-squared:	0.787
Method:	Least Squares	F-statistic:	371.4
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	2.82e-35
Time:	14:16:24	Log-Likelihood:	-57.048
No. Observations:	100	AIC:	116.1
Df Residuals:	99	BIC:	118.7
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
x1	0.3997	0.021	19.272	0.000	0.359	0.441
Omnibus:	0.720	Durbin-Watson:	2.036			
Prob(Omnibus):	0.698	Jarque-Bera (JB):	0.823			
Skew:	-0.105	Prob(JB):	0.663			
Kurtosis:	2.608	Cond. No.	1.00			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The estimated coefficient 0.4175 means with one unit increase in y, x will increase nearly 0.4175. Standard error rate shows the one standard deviation of β_y is $3.08e^{-8}$. So the 95% confidence interval of β_y is [0.383, 0.452] P-value close to 0 shows that y is statistically significant.

In [6]:

```
# part (c)
from IPython.display import Latex
Latex(r"""\begin{eqnarray}
\hat{x} = \hat{\beta}_x \times y \text{ ; versus ; } \hat{y} = \hat{\beta}_y \times x, \text{ the ; betas ;}
\text{ ; should ; be ; inverse ; of ; each ; other ; } (\hat{\beta}_x = \frac{1}{\hat{\beta}_y})
\text{ ; but ; they ; are ; somewhat ; off}

\end{eqnarray}""")
```

Out[6]:

```
\begin{eqnarray} \hat{x} = \hat{\beta}_x \times y \text{ ; versus ; } \hat{y} = \hat{\beta}_y \times x, \text{ the ; betas ;}
\text{ ; should ; be ; inverse ; of ; each ; other ; } (\hat{\beta}_x = \frac{1}{\hat{\beta}_y}) \text{ ; but ; they ; are ;}
\text{ ; somewhat ; off} \end{eqnarray}
```

part(e)

I used t- statistic formula from (d) and plug both (x,y) and (y,x) in it. I got same result.

In conclusion, the two regression lines should be the same just with the axes switched.

In [7]:

```
# part (f)

X = sm.add_constant(x)

model3 = sm.OLS(y,X).fit()

model3.summary()
```

Out[7]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.790
----------------	---	------------	-------

Dep. variable:	y	R-squared:	0.790
Model:	OLS	Adj. R-squared:	0.787
Method:	Least Squares	F-statistic:	367.6
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	6.22e-35
Time:	14:16:24	Log-Likelihood:	-136.94
No. Observations:	100	AIC:	277.9
Df Residuals:	98	BIC:	283.1
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	0.0042	0.096	0.044	0.965	-0.187	0.195
x1	1.9753	0.103	19.172	0.000	1.771	2.180

Omnibus:	0.555	Durbin-Watson:	1.775
Prob(Omnibus):	0.758	Jarque-Bera (JB):	0.687
Skew:	0.077	Prob(JB):	0.709
Kurtosis:	2.624	Cond. No.	1.07

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [8]:

```
Y = sm.add_constant(y)
model4 = sm.OLS(x,Y).fit()

model4.summary()
```

Out [8]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.790
Model:	OLS	Adj. R-squared:	0.787
Method:	Least Squares	F-statistic:	367.6
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	6.22e-35
Time:	14:16:24	Log-Likelihood:	-57.048
No. Observations:	100	AIC:	118.1
Df Residuals:	98	BIC:	123.3
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	0.0007	0.043	0.017	0.986	-0.085	0.087
x1	0.3997	0.021	19.172	0.000	0.358	0.441

Omnibus:	0.720	Durbin-Watson:	2.036
Prob(Omnibus):	0.698	Jarque-Bera (JB):	0.823
Skew:	-0.105	Prob(JB):	0.663
Kurtosis:	2.608	Cond. No.	2.07

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In conclusion, the t- statistic of two linear regression line are same.

exercise 3.12

In [9]:

```
# part(a)

from IPython.display import Latex
Latex(r"""\begin{eqnarray}
When \; x_{i}=y_{i}, or\; more\; generally\; when\; the \;beta \;denominators\; are\; equal\; \sum
x_{i}^2=\sum y_{i}^2

\end{eqnarray}""")
```

Out[9]:

```
\begin{eqnarray} When \; x_{i}=y_{i}, or\; more\; generally\; when\; the \;beta \;denominators\; are\; equal\; \sum x_{i}^2=\sum y_{i}^2
\end{eqnarray}
```

In [10]:

```
# part(b)

import pandas as pd

print("Model1's coefficient for Beta1 : ",model1.params)

print("Model2's coefficient for Beta2 : ",model2.params)
```

```
Model1's coefficient for Beta1 : [1.97540554]
Model2's coefficient for Beta2 : [0.39968647]
```

In [11]:

```
# part(c)
np.random.seed(111)

u = np.random.normal(1000,0.1,100)
v = np.random.normal(1000,0.1,100)

model5 = sm.OLS(v,u).fit()
model6 = sm.OLS(u,v).fit()

print("Model5's coefficient for Beta1 : ",model5.params)

print("Model6's coefficient for Beta2 : ",model6.params)
```

```
Model5's coefficient for Beta1 : [1.00000464]
Model6's coefficient for Beta2 : [0.99999534]
```

Excercise 3.13

In [12]:

```
#part (a)
np.random.seed(1111)

X1 = np.random.normal(0,1,100)
```

In [13]:

```
#part (b)
```

```
eps = np.random.normal(0,0.25,100)
```

In [14]:

```
#part(c)
Y1 = -1 + 0.5*X1 + eps

print(len(Y1))

Latex(r"""\begin{eqnarray}
\beta_{0}=-1 \\
\beta_{1}=0.5
\end{eqnarray}""")
```

100

Out[14]:

```
\begin{eqnarray} \beta_{0}=-1 \\ \beta_{1}=0.5 \end{eqnarray}
```


In [15]:

```
#part(d)
import matplotlib.pyplot as plt

plt.scatter(X1,Y1)
```

Out[15]:

<matplotlib.collections.PathCollection at 0x18136ff0cf8>



X and Y are nearly positively relative.

In [16]:

```
#part(e)
X2= sm.add_constant(X)

model8 = sm.OLS(Y1,X2).fit()

print(model8.params)

Latex(r"""\begin{eqnarray}
\beta_{0}= -0.75049222 \\
\beta_{1}= 0.44058306
\end{eqnarray}""")
```

[-1.04409137 0.01356952]


Out[16]:

```
\begin{eqnarray} \beta_{0}= -0.75049222 \\ \beta_{1}= 0.44058306 \end{eqnarray}
```

In [17]:

```
#part(f)
plt.scatter(X1,Y1)
plt.plot(X1,-1 + 0.5*X1,'r')
plt.plot(X1, -0.75049222 + 0.44058306*X1, 'b' )

plt.show()
```



In [18]:

```
#part(g)
import statsmodels.formula.api as smf

data = {"Y1": Y1, "X1": X1}

model9 = smf.ols(formula = 'Y1 ~ np.power(X1,2) + X1', data = data).fit()

print(model9.params)

table = sm.stats.anova_lm(model9,type = 2)

print(table)
```

```
Intercept          -0.996210
np.power(X1, 2)      0.005290
X1                  0.564663
dtype: float64
```

	df	sum_sq	mean_sq	F	PR(>F)
np.power(X1, 2)	1.0	0.009219	0.009219	0.134242	7.148715e-01
X1	1.0	27.351537	27.351537	398.277387	4.092429e-36
Residual	97.0	6.661435	0.068675	NaN	NaN

```
from sklearn.preprocessing import PolynomialFeatures from sklearn.linear_model import LinearRegression
```

```
poly_reg = PolynomialFeatures(2) xpoly = poly_reg.fit_transform(X1)
```

```
print(xpoly)
```

```
linearreg_2 = LinearRegression() linearreg_2.fit(xpoly,Y1)
```

```
plt.scatter(X1,Y1) plt.plot(X1,-1 + 0.5*X1,'r') plt.plot(X1, -0.75049222 + 0.44058306*X1, 'b' ) plt.plot(X1, linearreg_2.predict(xpoly))
```

```
plt.show()
```

In [19]:

```
#part(h)

np.random.seed(11111)

eps2 = np.random.normal(0,0.1,100)

Y2 = -1 + 0.5*X1 + eps2


model10 = sm.OLS(Y2,X2).fit()

print(model10.params)

plt.scatter(X1,Y2)
```

```
plt.plot(X1,-1 + 0.5*X1,'r')
plt.plot(X1, -0.7887987 + 0.09120311*X1, 'b' )
plt.show()
```

```
[-1.05679919  0.00845146]
```



In [20]:

```
# part(i)
np.random.seed(11111)

eps3 = np.random.normal(0,1,100)


Y3 = -1 + 0.5*X1 + eps3

model11 = sm.OLS(Y3,X2).fit()

print(model11.params)

plt.scatter(X1,Y3)
plt.plot(X1,-1 + 0.5*X1,'r')
plt.plot(X1, -0.45189999 + 0.07522537*X1, 'b' )
plt.show()
```

```
[-1.15037271  0.02118244]
```



part(j)

In [21]:

```
model8.summary()
```

Out [21]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.010
Method:	Least Squares	F-statistic:	0.04620
	Thu. 31 Jan	Prob (F-	

Date:	2019	statistic):	0.830
Time:	14:16:25	Log-Likelihood:	-87.962
No. Observations:	100	AIC:	179.9
Df Residuals:	98	BIC:	185.1
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0441	0.059	-17.723	0.000	-1.161	-0.927
x1	0.0136	0.063	0.215	0.830	-0.112	0.139

Omnibus:	0.419	Durbin-Watson:	1.794
Prob(Omnibus):	0.811	Jarque-Bera (JB):	0.572
Skew:	0.027	Prob(JB):	0.751
Kurtosis:	2.633	Cond. No.	1.07

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [22]:

```
model10.summary()
```

Out [22]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.010
Method:	Least Squares	F-statistic:	0.02733
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	0.869
Time:	14:16:25	Log-Likelihood:	-66.854
No. Observations:	100	AIC:	137.7
Df Residuals:	98	BIC:	142.9
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0568	0.048	-22.155	0.000	-1.151	-0.962
x1	0.0085	0.051	0.165	0.869	-0.093	0.110

Omnibus:	0.074	Durbin-Watson:	1.959
Prob(Omnibus):	0.964	Jarque-Bera (JB):	0.235
Skew:	0.029	Prob(JB):	0.889
Kurtosis:	2.770	Cond. No.	1.07

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [23]:

```
model11.summary()
```

Out [23]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.010
Method:	Least Squares	F-statistic:	0.03133
Date:	Thu, 31 Jan 2019	Prob (F-statistic):	0.860
Time:	14:16:25	Log-Likelihood:	-151.91
No. Observations:	100	AIC:	307.8
Df Residuals:	98	BIC:	313.0
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.1504	0.112	-10.301	0.000	-1.372	-0.929
x1	0.0212	0.120	0.177	0.860	-0.216	0.259

Omnibus:	2.211	Durbin-Watson:	1.959
Prob(Omnibus):	0.331	Jarque-Bera (JB):	1.785
Skew:	-0.176	Prob(JB):	0.410
Kurtosis:	2.448	Cond. No.	1.07

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [24]:

```
from matplotlib import pyplot as plt
from pandas.tools.plotting import scatter_matrix
from mpl_toolkits.mplot3d import Axes3D
from statsmodels.stats.outliers_influence import OLSInfluence

%matplotlib inline
plt.style.use('ggplot')
```

Exercise 3.14

In [25]:

```
#part(a)

np.random.seed(2)

a1 = np.random.random(100)

a2 = 0.5*a1 + np.random.randn(100)/10

b = 2+2* a1 +0.3* a2 + np.random.randn(100)

data2 = {"b":b , "a1": a1, "a2": a2}
```

In [26]:

```
#part(b)
from numpy import corrcoef
print(corrcoef(a1,a2))


df = pd.DataFrame(np.column_stack((b,a1,a2)), columns=['b','a1','a2'] )
df.head()

scatter_matrix(df, figsize = (8,8),alpha=1);
```

```
[[1.          0.81145744]
 [0.81145744 1.          ]]
```

C:\Users\husiw\Anaconda3\lib\site-packages\ipykernel_launcher.py:9: FutureWarning: 'pandas.tools.plotting.scatter_matrix' is deprecated, import 'pandas.plotting.scatter_matrix' instead.

```
if __name__ == '__main__':
```



In [27]:

```
#part(c)
modell12 = smf.ols(formula = 'b ~ a1 + a2', data = data2).fit()

print(modell12.summary())

print(modell12.params)
```

OLS Regression Results

```
=====
Dep. Variable:          b    R-squared:                0.335
Model:                  OLS    Adj. R-squared:          0.322
Method:                 Least Squares    F-statistic:      24.47
Date:                   Thu, 31 Jan 2019    Prob (F-statistic): 2.48e-09
Time:                   14:16:26    Log-Likelihood:    -137.23
No. Observations:       100    AIC:               280.5
Df Residuals:           97    BIC:               288.3
Df Model:                2
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.7199	0.195	8.817	0.000	1.333	2.107
a1	1.8561	0.624	2.975	0.004	0.618	3.095
a2	1.1244	0.874	1.287	0.201	-0.610	2.859

```
=====
Omnibus:                 0.358    Durbin-Watson:          1.843
Prob(Omnibus):           0.836    Jarque-Bera (JB):        0.501
Skew:                    -0.120    Prob(JB):                0.778
Kurtosis:                 2.749    Cond. No.                 12.1
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

Intercept    1.719888
a1           1.856144
a2           1.124409
dtype: float64

```

The intercept $\beta_0 = 1.72$ is significant and β_1 is barely significant so we reject the hypothesis that $\beta_0 = 0$ and the hypothesis that $\beta_1 = 0$ but we can't reject the hypothesis that $\beta_2 = 0$. Also notice the SE and confidence intervals for all three coefficients are very large.

In [28]:

```

# part(c)-(e)
modell3 = smf.ols(formula = 'b ~ a1', data = data2).fit()

print(modell3.summary())

modell4 = smf.ols(formula = 'b ~ a2', data = data2).fit()

print(modell4.summary())

```

OLS Regression Results

```

=====
Dep. Variable:          b    R-squared:                0.324
Model:                  OLS    Adj. R-squared:          0.317
Method:                 Least Squares    F-statistic:        46.98
Date:                   Thu, 31 Jan 2019    Prob (F-statistic):    6.43e-10
Time:                   14:16:26    Log-Likelihood:        -138.08
No. Observations:       100    AIC:                   280.2
Df Residuals:           98    BIC:                   285.4
Df Model:               1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.6908	0.194	8.698	0.000	1.305	2.077
a1	2.5078	0.366	6.854	0.000	1.782	3.234

```

=====
Omnibus:                 0.247    Durbin-Watson:          1.868
Prob(Omnibus):           0.884    Jarque-Bera (JB):        0.341
Skew:                    -0.112    Prob(JB):                0.843
Kurtosis:                 2.823    Cond. No.:               4.61
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```

=====
Dep. Variable:          b    R-squared:                0.275
Model:                  OLS    Adj. R-squared:          0.267
Method:                 Least Squares    F-statistic:        37.12
Date:                   Thu, 31 Jan 2019    Prob (F-statistic):    2.18e-08
Time:                   14:16:26    Log-Likelihood:        -141.60
No. Observations:       100    AIC:                   287.2
Df Residuals:           98    BIC:                   292.4
Df Model:               1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.0661	0.163	12.702	0.000	1.743	2.389
a2	3.2334	0.531	6.093	0.000	2.180	4.286

```

=====
Omnibus:                 1.173    Durbin-Watson:          1.787
Prob(Omnibus):           0.556    Jarque-Bera (JB):        1.095
Skew:                    -0.093    Prob(JB):                0.578
Kurtosis:                 2.522    Cond. No.:               5.58
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

part(f)

In model13, β_1 , we find that the coefficient is close to the true value and is now very significant. In model14 β_2 , we find that the coefficient for a_2 is now very significant. Note there is no contradiction here. a_1 and a_2 are strongly correlated and each is related to b independent of each other.

In [31]:

```
df.loc[len(df)] = [0.1, 0.8, 6]

A = sm.add_constant(df[['a1', 'a2']])
model = sm.OLS(df.b, A)
estimate = model.fit()

print(estimate.summary())

# Obtain the residuals, studentized residuals and the leverages
fitted_values = estimate.fittedvalues.values
residuals = estimate.resid.values
studentized_residuals = OLSInfluence(estimate).resid_studentized_internal
leverages = OLSInfluence(estimate).influence

# Plot
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 4))

# Studentized Residuals
ax1.scatter(fitted_values[:-1], studentized_residuals[:-1], facecolors='none', edgecolors='b');
# Plot the possible Outlier in red
ax1.scatter(fitted_values[-1], studentized_residuals[-1], facecolors='none', edgecolors='r');
ax1.set_xlabel('fitted values');
ax1.set_ylabel('studentized residuals');

# Leverages
ax2.scatter(leverages[:-1], studentized_residuals[:-1], facecolors='none', edgecolors='b');
# plot the possible high leverager in red
ax2.scatter(leverages[-1], studentized_residuals[-1], facecolors='none', edgecolors='r');
ax2.set_xlabel('Leverage');
ax2.set_ylabel('studentized residual');
```

OLS Regression Results

```
=====
Dep. Variable:          b      R-squared:                0.403
Model:                  OLS      Adj. R-squared:          0.391
Method:                 Least Squares      F-statistic:        33.77
Date:                   Thu, 31 Jan 2019      Prob (F-statistic):    6.22e-12
Time:                   14:16:26      Log-Likelihood:       -141.94
No. Observations:       103      AIC:                 289.9
Df Residuals:           100      BIC:                 297.8
Df Model:                2
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	1.6794	0.195	8.611	0.000	1.292	2.066
a1	2.8495	0.384	7.423	0.000	2.088	3.611
a2	-0.6217	0.104	-5.961	0.000	-0.829	-0.415

```
=====
Omnibus:                 0.198      Durbin-Watson:          1.908
Prob(Omnibus):           0.906      Jarque-Bera (JB):        0.196
Skew:                    -0.098      Prob(JB):                0.907
Kurtosis:                2.913      Cond. No.:               5.70
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

